


Image Segmentation: Edge-based

Michael A. Wirth, Ph.D.
University of Guelph
Computing and Information Science
Image Processing Group
© 2004

Edge-based Segmentation

- The history of edge detection



Laplacian of Gaussian

- The Laplacian is seldom used on its own for edge detection because of its sensitivity to noise.
- The **Laplacian-of-Gaussian** (LoG) uses a Gaussian filter to blur the image and a Laplacian to enhance edges.
 - Also known as **Marr & Hildreth** edge detector
- Edge localisation is done by finding **zero-crossings**.

Laplacian of Gaussian

- Reference:
 - Marr, D., and Hildreth, E., “Theory of edge detection”, *Proceedings of the Royal Society of London, Series B*, 1980, **207**:pp.187-217

Laplacian of Gaussian

Algorithm: LoG

- Convolve the image with a two-dimensional Gaussian function
- Compute the Laplacian of the convolved image $\rightarrow L$
- Identify edge pixels as those for which there is a zero-crossing in L .

Laplacian of Gaussian

- A radially-symmetric 2D Gaussian:

$$G_{\theta}(x, y) = e^{\left(\frac{-r^2}{2\sigma^2}\right)}$$


where $r^2 = x^2 + y^2$

- The Laplacian of this is:

$$G(x, y) = \left(\frac{r^2 - \sigma^2}{\sigma^4}\right) e^{\left(-\frac{r^2}{2\sigma^2}\right)}$$

Laplacian of Gaussian

- This function has a minimum at its origin, but it is usual to invert the filter, so that it has a maximum at its origin.
 - classic “Mexican hat” shape



Laplacian of Gaussian

- The value of σ determines the width of the filter and controls the amount of smoothing produced by the Gaussian component.
 - σ tunes the filter to detect edges at different scales
 - Should have a half-width of at least 3σ

Laplacian of Gaussian

- 5×5 Laplacian of Gaussian kernel

$$\begin{bmatrix} 0 & 0 & -1 & 0 & 0 \\ 0 & -1 & -2 & -1 & 0 \\ -1 & -2 & 16 & -2 & -1 \\ 0 & -1 & -2 & -1 & 0 \\ 0 & 0 & -1 & 0 & 0 \end{bmatrix}$$

Difference of Gaussian

- The LoG can be approximated by convolving with a kernel that is the difference of two Gaussian kernels with substantially different σ 's.
 - Known as the **Difference-of-Gaussian** (DoG)

Zero-Crossings

- To form an edge map from the LoG/DoG output → locate and mark the zero-crossings
 - Trying to detect zeros in the LoG or DoG image will fail
 - A simple zero-crossing detector may identify a zero-crossing in a $n \times n$ window, assigning an edge label if LoG/DoG image values of both polarities occur in the window
 - No edge label would be assigned if values within the window are all of one sign (either all positive, or all negative)



Scale Space

- Gaussian smoothing results in the blurring of edges and other sharp discontinuities in an image.
 - The amount of smoothing depends on the value of σ .
 - A larger σ results in better noise filtering, but at the same time loses important edge information, affecting the performance of the edge detector.
 - If a small filter is used, there is likely to be more noise due to insufficient averaging.

Scale Space



- In order to ensure a variety of scales are used:
 - Use two different Gaussians, say $\sigma \pm 0.8$
e.g. if $\sigma=2.0$, then $\sigma_1=1.2$, $\sigma_2=2.8$
 - Select the pixels in the edge images that have zero-crossings in both scales as the edge pixels
 - More than two Gaussians can be used.

Examples of LoG



$\sigma=1.5$


$\sigma=2.3$




$\sigma=3.0$

Examples of LoG

$\sigma=1.0$




$\sigma=1.4$



$\sigma=1.5$



$\sigma=2.0$



Limitations of 2nd Derivatives

- It smoothes the shape too much
- It tends to create closed loops for edges

Extensions to Zero-Crossing

- Qian, R.J., and Huang, T.S., “Optimal edge detection in two-dimensional images”, in *ARPA Image Understanding Workshop*, Monterey, CA, ARPA, 1994, pp.1581-1588
- Mehrotra, R., and Shiming, Z., “A computational approach to zero-crossing-based two-dimensional edge detection”, *Graphical Models and Image Processing*, 1996, **58**:pp.1-17
- Hardie, R.C. and Boncelet, C.G., “Gradient-based edge detection using nonlinear edge-enhancing filters”, *IEEE Transactions on Image Processing*, 1995, **4**:1572-1577
- Alparone, L., Baronti, S., and Casini, A., “A novel approach to the suppression of false contours”, in *International Conference on Image Processing*, IEEE, 1996, pp.825-828

Canny Edge Detector

- The **Canny** edge detector addresses the fact that for edge detection, there is a tradeoff between noise reduction (smoothing) and edge localisation.
 - A form of optimal edge detection
- Reference:
 - Canny, J., “A computational approach to edge detection”, *IEEE Transactions on Pattern Analysis and Machine Intelligence*, 1986, **8**(6):pp.679-698

Canny Edge Detector

Algorithm: Canny Edge Detection

- Smooth the image with a Gaussian filter
- Compute the gradient magnitude and orientation
- Apply *non-maximal suppression* to the gradient magnitude image
- Use *hysteresis thresholding* to detect and link edges

Canny Edge Detector

- The **smoothing** step employs a Gaussian low-pass filter:
 - The standard deviation, σ , determines the width of the filter and hence the amount of smoothing
 - A filter with a large σ will suppress much of the noise, but also smooth away the weakest edges

Canny Edge Detector

- The **edge enhancement** step simply involves calculation of the gradient vector at each pixel in the smoothed image.
- Efficient implementations combine the smoothing and enhancement steps by convolving the image with a derivative of the Gaussian kernel.

Canny Edge Detector

- The **localization** step has two stages:
 - Non-maximal suppression
 - Hysteresis thresholding

Canny Edge Detector

- Non-maximal suppression:
 - Pixels that are not local maxima are removed
 - Thins the wide ridges around local maxima in gradient magnitude images down to edges that are only one pixel wide

Canny Edge Detector

Algorithm: Non-maximal Suppression

- Quantize edge directions eight ways according to 8-connectivity
- For each pixel with non-zero edge magnitude, inspect the two adjacent pixels indicated by the direction of its edge
- If the edge magnitude of either of these two exceeds that of the pixel under inspection, mark it for deletion
- When all pixels have been inspected, erase to zero all edge pixels marked for deletion

Canny Edge Detector

- Hysteresis thresholding:
 - Assumes the gradient magnitude G_M and direction G_θ have already been computed.
 - There are problems associated with applying a single, fixed threshold to gradient magnitude images.
 - Choosing a low threshold ensures that weak, yet meaningful edges are captured, but may also result in an excessive number of “false-positives”.
 - Gradient maxima are caused by noise rather than by features of interest.
 - Choosing a high threshold will lead to excessive fragmentation of pixels that represent significant edges in the image

Canny Edge Detector



Algorithm: Hysteresis thresholding

- Uses two thresholds T_{low} and T_{high}
- If an edge response is above T_{high} , those pixels constitute a definite edge. Individual weak responses usually correspond to noise.
- If weak responses are connected to pixels with strong responses and they are above T_{low} they are edge pixels
- The T_{low} and T_{high} thresholds are set according to an estimated signal-to-noise ratio

Canny Edge Detector



- In some versions the search is conducted over all eight-neighbors of an edge pixel (edgel), in others, only the neighbors along a line normal to the gradient orientation at the edge pixel are considered.
 - This technique reduces the number of false-positives because edges are tracked only if at least one pixel has a gradient magnitude exceeding T_{high}
 - It reduces the fragmentation of contours in the edge map by allowing significant fluctuations to occur in gradient magnitude on an edge.

Examples of Canny



$\sigma=0.6$

$\sigma=1.5$



$\sigma=2.3$

Extensions to Canny

- Jalali, S., and Boyce, J.F., “Determination of optimal general edge detectors by global minimization of a cost function”, *Image and Vision Computing*, 1995, **13**:pp.683-693
- Sorrenti, D.G., “A proposal on local and adaptive determination of filter scale for edge detection”, in *Image Analysis and Processing, ICIAP’95*, Springer Verlag, 1995, pp.405-410
- Demigny, D., Lorca, F.G., Kessal, L., “Evaluation of edge detector performances with a discrete expression of Canny’s criteria”, in *International Conference on Image Processing*, IEEE, 1995, pp.169-172
- Mehrotra, R., and Shiming, Z., “A computational approach to zero-crossing-based two-dimensional edge detection”, *Graphical Models and Image Processing*, 1996, **58**:pp.1-17

Shen-Castan Detector

- The **Shen-Castan** edge detector uses an optima filter function:
 - **Infinite Symmetric Exponential Filter** (ISEF)
 - Better signal-to-noise ratios and better localisation than Canny
- Reference:
 - Shen, J., and Castan, S., “An optimal linear operator for step edge detection”, *CVGIP: Graphical Models and Understanding*, 1992, **54**(2):pp.112-133

Shen-Castan Detector

Algorithm: Shen-Castan

- Convolve the image using the ISEF
- Localise edges by finding zero-crossings of the Laplacian (similar to the Marr-Hildreth algorithm)
 - Approximate the Laplacian by subtracting the original from the smoothed image, the resulting image is a *band-limited Laplacian*
 - Compute the *Binary Laplacian Image* (BLI) by setting all the positive valued pixels to 1, and others to 0.
 - The candidate pixels are on the boundaries of the regions in the BLI


Shen-Castan Detector

- Improve the quality of the edge pixels by *false zero-crossing suppression*
 - At the location of an edge pixel there will be a zero-crossing in the 2nd derivative of the filtered image. This means the gradient at that point is either a maximum or a minimum
 - If the change is +ve to -ve → positive zero-crossing
 - If the change is -ve to +ve → negative zero-crossing
 - Allow +ve zero-crossing to have a +ve gradient, and -ve zero-crossings to have a -ve gradient
 - Suppress all other zero-crossings

Shen-Castan Detector

- Perform adaptive gradient thresholding
 - A window with fixed width W is centred on the candidate edge pixel. If this is an edge pixel then the window will contain two regions of differing intensity separated by an edge (zero crossing contour).
 - The best estimate of the gradient at that point should be the difference in level between the two regions
- Apply hysteresis thresholding

Example of Shen-Castan



% of pixels above the
hysteresis threshold = 0.7
smoothing factor = 0.7
window size = 15
thinning factor
(number of pixels) = 1

Edge Image Segmentation

- An edge image is usually thresholded to decide which edges are significant, and **streaking** is the breaking up of an edge contour caused by the operator fluctuating above and below the threshold.
 - Can be eliminated by thresholding using some form of edge image segmentation

Edge Image Segmentation

- Edge image thresholding
 - Non-maximal suppression, and hysteresis
- Edge relaxation
- Border tracing
- Border detection as graph searching
- Border detection as dynamic programming

Other Edge Segmentation Techniques

- Boie & Cox Edge Detector
 - Boie, R.A. Cox, I., and Rehak, P., “On optimum edge recognition using matched filters”, in *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition*, 1986, pp.100-108
 - Boie, R.A. and Cox, I., “Two-dimensional optimum edge recognition using matched and wiener filters for machine vision”, in *Proceedings of the IEEE First International Conference on Computer Vision*, 1987, pp.450-456

Image Segmentation: Region-based

Michael A. Wirth, Ph.D.
University of Guelph
Computing and Information Science
Image Processing Group
© 2004

Pixel Connectivity

- Pixel connectivity is a central concept in both edge- and region-based approaches to segmentation.
- Define two types of neighborhood surrounding a pixel:
 - **4-neighborhood**: contains pixels, above, below, to the left and to the right of the central pixel.
 - **8-neighborhood**: contains all the pixels of a 4-neighborhood, plus four diagonal neighbors.

Pixel Connectivity

- A 4-connected path from pixel p_1 to another pixel p_n is the sequences of pixels $\{p_1, p_2, \dots, p_n\}$, where p_{i+1} is a 4-neighbor of p_i for all $i=1, \dots, n-1$
- The path is said to be 8-connected if p_{i+1} is an 8-neighbor of p_i

Region Segmentation

- Given a set of image pixels I , find a partition S of the image I into a set of n regions R_i :

$$\bigcup_{i=1}^n R_i = I$$

Region Similarity

- The **uniformity** or otherwise of a connected region of pixels may be indicated by a **uniformity predicate** \rightarrow a logical statement that is true only if pixels in the region are sufficiently similar in terms of intensity, colour, or some other property.
- A common uniformity predicate:

$$P(R) = \begin{cases} TRUE & \text{if } |f(i, j) - f(k, l)| \leq \Delta \\ FALSE & \text{otherwise} \end{cases}$$

Region Similarity

- Where (i,j) and (k,l) are the coordinates of neighbouring pixels in region R .
- It states that a region, R is uniform if (and only if) any two neighboring pixels differ in intensity by no more than Δ .
- Small changes in intensity from neighbour to neighbour can accumulate resulting in a large difference in intensity between opposite sides of a region.

Region Similarity

- A similar predicate:

$$P(R) = \begin{cases} TRUE & \text{if } |f(i,j) - \mu_R| \leq \Delta \\ FALSE & \text{otherwise} \end{cases}$$

- Where $f(i,j)$ is the intensity value of a pixel from region R with coordinate (i,j) , and μ_R is the mean intensity of all pixels in R except the pixel at (i,j) .

Region Growing

- **Region growing** is bottom-up algorithm that starts with a group of pixels called **seeds** that belong to the structure of interest.
 - The aim is to grow a uniform, connected region from each seed.
 - Comparing the difference between the pixel intensity value and the mean intensity value over a region.
 - The results of region growing depend strongly on the selection of the homogeneity criterion.
 - If it is not properly chosen, the regions **leak** out into adjoining areas or merge with regions that do not belong to the structure of interest.
 - Different starting points may not grow into identical regions.

Region Growing

- A pixel is included in a region if:
 - It has not been assigned to any other region
 - It is a neighbour of that region
 - The new region created by addition of the pixel is still uniform.

Region Growing

- Segmentation of an image must satisfy a number of criteria:
 - ❶ all pixels must be assigned to regions
 - ❷ each pixel must belong to a single region only
 - ❸ each region must be a connected set of pixels
 - ❹ each region must be uniform
 - ❺ any merged pair of adjacent regions must be nonuniform

Region Growing

- Region growing satisfies criteria ③ and ④
- Criteria ① and ② are not satisfied:
 - In general the number of seeds defined will not be sufficient to create a region for every pixel.
- Criteria ⑤:
 - Regions grown from two nearby seeds are always regarded as distinct even if those seeds are defined in a part of the image that should be segmented as a single region

Region Growing

- Various approaches to region growing:
 - Zucker, S.W., “Region growing: Childhood and adolescence”, *Computer Graphics and Image Processing*, 1976, **5**, pp.382-399
 - Fu, K.S., and Mui, J.K., “A survey on image segmentation”, *Pattern Recognition*, 1981, **13**(1):pp.3-16
 - Haralick, R.M., and Shapiro, L.G., “Survey: Image segmentation techniques”, *Computer Vision, Graphics and Image Processing*, 1985, **29**:pp.100-132

Split and Merge

- Complete segmentation is possible if a top-down approach is adopted.
 - The entire region is considered initially to be a single region
 - Divide the region into subregions
 - These subregions are then split or merged in an attempt to meet the uniformity criteria.
 - The algorithm iterates until all regions are uniform or until the desired number of regions have been established:

Algorithm

Algorithm: Split and Merge

1. Start with the entire image as a single region
2. Choose a region R . If $P(R)$ is false, split the region into four subregions
3. Consider any two or more neighbouring subregions, R_1, R_2, \dots, R_n , in the image. If $P(R_1 \cup R_2 \cup \dots \cup R_n)$ is true, merge the n regions into a single region.
4. Repeat these steps until no further splits or merges take place.

Region Splitting

- If region splitting alone were used to segment the image, the final partition would likely contain many small, adjacent regions with identical properties.
- Alternate region splitting with a merging stage in which two adjacent regions R_i and R_j are combined into a new larger region if the uniformity predicate for the union of these regions is true.

$$P(R_i \cup R_j) \text{ is } TRUE$$

Algorithm

Algorithm: Region Splitting

1. Form initial regions in the image
2. For each region in an image, perform the following steps:
 - (a) Compute the variance in the intensity value for the region
 - (b) If the variance is above a threshold, split the region
3. Repeat step 2 until no regions are split

Region Merging

- Determine the similarity between two regions.
 - Approaches are based on the intensity of regions or on the weakness of boundaries between the regions.
- There are two approaches to judging the similarity of adjacent regions:
 - Compare their mean intensities. If the mean intensities do not differ by more than some predetermined values, the regions are considered similar and should be candidates for merging.

Region Merging

- Assume that the intensity values are drawn from a probability distribution. Consider whether or not to merge adjacent regions based on the probability that they will have the same statistical distribution of intensity values.

Algorithm

Algorithm: Region Merging

1. Form initial regions in the image
2. For each region in an image, perform the following steps:
 - (a) Consider its adjacent region and test to see if they are similar
 - (b) If the regions are similar, merge them
3. Repeat step 2 until no regions are merged

Region Labelling

- Find a connected region of pixels that were detected by thresholding, and give all the pixels in that region their own unique value.
- The simplest and most common labeling algorithm scans the image pixel by pixel invoking a recursive labeling procedure whenever a non-zero pixel is found.

Region Labelling

- The algorithm implements the “grassfire” concept
 - Imagine a “fire” is started at a pixel, and it propagates to any of the pixel’s 4- or 8- neighbours that were also detected by thresholding.
 - Wherever a fire is started, the pixel is “burnt-away” (has its value set to zero), so that it cannot be visited again by the labelling algorithm.
 - At the end of the algorithm all pixels belonging to the region have been set to zero in the input image → making it indistinguishable from the background, and the corresponding pixels in the output image have been assigned a region number.
 - The region number is then incremented, ready for the next connected region.